

Retrieval Augmented Generation for Enterprise Applications 2025

Sources: 13 | Findings: 18 | Tokens: 4,200 | Cost: \$0.0630 | Generated: 2026-04-16T13:26:11.809812+00:00

Executive Summary

Retrieval-Augmented Generation (RAG) has become the dominant paradigm for grounding large language model outputs in verified enterprise knowledge. Organisations adopting RAG report measurable reductions in hallucination rates and improved compliance with internal documentation standards.

The technology is now considered production-ready for knowledge management, customer support, and regulatory reporting use cases across multiple industries.

<https://arxiv.org/abs/2312.10997>

<https://example.com/enterprise-rag-2025>

Key Findings

Hybrid search combining dense vector retrieval with BM25 keyword matching consistently outperforms either approach in isolation across enterprise benchmarks.

Domain-specific embedding models fine-tuned on proprietary corpora deliver a 15-30% improvement in retrieval precision compared to general-purpose models. Reranking retrieved chunks prior to generation is now considered a standard best practice for production deployments.

<https://arxiv.org/abs/2401.00368>